

## Part 1

### Stack Frame

%rbp	Saved base pointer
-1(%rbp)	Return value (8-bit)
-2(%rbp)	
-3(%rbp)	
-4(%rbp)	
-5(%rbp)	Current (32-bit)
-6(%rbp)	
-7(%rbp)	
-8(%rbp)	
-9(%rbp)	
-10(%rbp)	
-11(%rbp)	
-12(%rbp)	
-13(%rbp)	
-14(%rbp)	
%rsp	

All other values are initially accessed through the argument registers (except for array, which is located in the stack and explained in part 2), and assigned variable names which are then used to access those values.

The reason for the number of unused bytes is that x86-64 requires stack frames to be a multiple of 16 bytes, even though the function does not necessarily need that much memory.

The return value is a Boolean so only 8 bits are used to store it, while current is an integer so 32 bits are used instead.

Variables other than current are initialised in badrandom and referenced by name.

Recursion here is inefficient in terms of memory because a new stack frame is created with every recursive call, reserving another 16 bytes of memory each time.

Also note that %r14, %r15 and %rbx do not need to be saved in the unoptimised code because they are never used.

## Part 2

### Stack frame

%rbp	Saved base pointer
-1(%rbp)	
-2(%rbp)	
-3(%rbp)	
-4(%rbp)	
-5(%rbp)	
-6(%rbp)	
-7(%rbp)	Saved value of %r15 (64-bit)
-8(%rbp)	
-9(%rbp)	Saved value of %r14 (64-bit)
-10(%rbp)	
-11(%rbp)	
-12(%rbp)	
-13(%rbp)	
-14(%rbp)	
-15(%rbp)	
-16(%rbp)	
-17(%rbp)	Saved value of %rbx
-18(%rbp)	
-19(%rbp)	
-20(%rbp)	
-21(%rbp)	
-22(%rbp)	
-23(%rbp)	
-24(%rbp)	

The optimised assembly code never explicitly reserves memory space for the stack frame. Instead, the only space taken up on the stack is by pushing callee-saved values in indexes %r14, %r15 and %rbx, as well as for saving the previous base pointer. This is to follow the function call convention of preserving caller-owned registers if the callee wants to use them. The other values, which are all initialised in badrandom, are mostly stored in dedicated registers. The exception is the array pointer, which is in the stack above the base pointer for this function, and its address in the caller's stack frame is accessed using 16(%rbp) when needed (there are only 6 dedicated argument registers, and \*in\_array is the 7<sup>th</sup> argument to badrandom).

Start/current (32-bit)	%edi
Mult/in_mult	%esi
Add/in_add	%edx
Div/in_div	%ecx
Min/in_min	%r8d
*in_length	%r9
*array/*in_array	16(%rbp)

## Optimisations

The variables `*array`, `mult`, `add`, `div` and `min` do not take up space in the stack frame. Instead, as they are all initialised to arguments to `badrandom`, the argument indexes (%edi to %r9) are used instead. While upon first glance this might seem problematic, as single registers are being used to track the values of two separate variables, it works out fine as the arguments are never needed again after being used to initialise the variables. For example, when `badrandom_recurse(start)` is called, `current` is set to the same value as `start`. After this point, the computer no longer uses the `start` variable, so instead treats the corresponding register (%edi) as storing the value of `current`. Note that this also eliminates the need for these values to be assigned variable names.

Another optimisation made is that this assembly code avoids recursion. While `badrandom_recurse` is written recursively in the C code and is treated as such in the unoptimised assembly, it is now treated as an iterative function. The `callq` instruction is never used to make a recursive call and additional space is never reserved in the stack frame. Instead, the assembly repeatedly jumps back to the `.LBB0_3` label every time a recursive call is made, and runs the same operations on the new value of `current`. On each iteration, it checks the two termination conditions and exits the loop whenever one of them is met. This is why it never makes a call to `badrandom_recurse`, as it is effectively treating the whole process as a single function with an iterative loop inside of it.

The key difference between the `call` and `jump` instruction is that `jump` simply moves to a given label, while `call` pushes the current location it will return to in the stack before moving the instruction pointer. This means that each `call` uses up stack space while `jumping` does not.

This code also simplifies the handling of the `if/else` recursive call in `badrandom_recurse` (shown below for reference). The only difference between the `if` and `else` blocks is the sign of `add`. It is positive in the `if` block and negative in the `else` block. The unoptimised code compiles the below code into 3 sections, for the condition checking, `if` block and `else` block, which seems like a fairly straightforward way of handling it. However, the optimised assembly code recognises that the blocks are doing the exact same thing with slightly different values, and stores both `+add` and `-add` in separate registers to be used in the operation depending on what the `if` condition evaluates to.

```
if (current%min == 0) {
    return badrandom_recurse((current*mult+add)%div); }
else {
    return badrandom_recurse((current*mult-add)%div); }
```

}
---